

## MNIST Quick Overview:

Data: 70K handwritten digits (60K train, 10K test), each 28x28 grayscale

Model: Classifier neural network (10 output classes, one per digit. Model predicts probability of digit for each example.

Use Cases: Basic digit recognition (postal codes, checks, forms)

Why Used: Perfect for learning ML/testing ideas (small, clean dataset)

Limitation: Too simple for real applications (only digits, clean images)

```
In [88]: from torchvision import datasets
from torchvision.transforms import ToTensor

# Download MNIST
train_data = datasets.MNIST(
    root = 'data',
    train = True,
    transform = ToTensor(),
    download = True
)

test_data = datasets.MNIST(
    root = 'data',
    train = False,
    transform = ToTensor()
)

# Convert to numpy arrays
x_train = train_data.data.numpy()
y_train = train_data.targets.numpy()
x_test = test_data.data.numpy()
y_test = test_data.targets.numpy()

# Reshape and normalize if needed
x_train = x_train.astype('float32') / 255
x_test = x_test.astype('float32') / 255

print(f"x_train shape: {x_train.shape}") # Should be (60000, 28, 28)
print(f"x_test shape: {x_test.shape}") # Should be (10000, 28, 28)
print(f"y_train shape: {y_train.shape}") # Should be (60000,)
print(f"y_test shape: {y_test.shape}") # Should be (10000,)

x_train shape: (60000, 28, 28)
x_test shape: (10000, 28, 28)
y_train shape: (60000,)
y_test shape: (10000,)
```

```
In [90]: # Data prep for training

x_train_scaled = torch.FloatTensor(x_train)
x_test_scaled = torch.FloatTensor(x_test)

y_train = torch.FloatTensor(y_train.tolist()).long()
y_test = torch.FloatTensor(y_test.tolist()).long()
```

```
In [91]: y_train
```

```
Out[91]: tensor([5, 0, 4, ..., 5, 6, 8])
```

```
In [81]: import torch
import torch.nn as nn

class ClassificationNet(nn.Module):
    def __init__(self, num_classes=10):
        super(ClassificationNet, self).__init__()
        # First flatten the 28x28 input to 784 (28*28)
        self.flatten = nn.Flatten()

        # Layer 1
        self.layer_1 = nn.Linear(784, 512) # 784 = 28*28
        self.bn1 = nn.BatchNorm1d(512)

        # Layer 2
        self.layer_2 = nn.Linear(512, 256)
        self.bn2 = nn.BatchNorm1d(256)

        # Layer 3
        self.layer_3 = nn.Linear(256, num_classes)

        self.relu = nn.ReLU()
```

```

        self.dropout = nn.Dropout(0.2)

    def forward(self, x):
        # Flatten the input
        x = self.flatten(x)

        # Forward pass through layers
        x = self.dropout(self.relu(self.bn1(self.layer_1(x))))
        x = self.dropout(self.relu(self.bn2(self.layer_2(x))))
        x = self.layer_3(x)
        return x

```

In [ ]: # Model training

```

model = ClassificationNet()
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr = 0.001)
num_epoch = 1000
train_loss_plot = []
loss_test = []
for epoch in range(num_epoch):
    model.train()
    outputs = model(x_train_scaled)
    loss = criterion(outputs, y_train)

    optimizer.zero_grad()
    loss.backward()
    optimizer.step()

    # Test accuracy
    train_loss_plot.append(loss.item())
    loss_test.append(criterion(model(x_test_scaled), y_test).item())

```


In [92]: # Model training evaluate

```

import matplotlib.pyplot as plt
plt.figure(figsize=(10, 5))
plt.subplot(1, 2, 1)
plt.plot(train_loss_plot, label='Training Loss')
plt.plot(loss_test, label='Val Loss')
plt.title('Training Val Loss Over Time')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.grid(True)
plt.legend(loc='upper right') # Specify location

```

Out [92]: <matplotlib.legend.Legend at 0x1b2a05550>



In [84]: from sklearn.metrics import precision\_recall\_fscore\_support, classification\_report

```

# Evaluate on unseen test data
model.eval()
with torch.no_grad():
    # Get predictions

```

```

test_outputs = model(x_test_scaled)

# Calculate loss
test_loss = criterion(test_outputs, y_test)

# Get predicted classes
_, predicted = torch.max(test_outputs.data, 1)

# Convert tensors to numpy for sklearn metrics
y_true = y_test.cpu().numpy()
y_pred = predicted.cpu().numpy()

# Calculate accuracy
total = y_test.size(0)
correct = (predicted == y_test).sum().item()
accuracy = 100 * correct / total

# Calculate precision and recall for each class
precision, recall, f1, _ = precision_recall_fscore_support(y_true, y_pred, average=None)

# Print metrics
print(f'Test Loss: {test_loss.item():.4f}')
print(f'Test Accuracy: {accuracy:.2f}%')

print("\nMetrics per class:")
for i in range(len(precision)):
    print(f"\nClass {i}:")
    print(f"Precision: {precision[i]:.4f}")
    print(f"Recall: {recall[i]:.4f}")
    print(f"F1-score: {f1[i]:.4f}")

# Print detailed classification report
print("\nDetailed Classification Report:")
print(classification_report(y_true, y_pred))

```

Test Loss: 1.4937  
Test Accuracy: 97.52%

Metrics per class:

Class 0:  
Precision: 0.9798  
Recall: 0.9908  
F1-score: 0.9853

Class 1:  
Precision: 0.9826  
Recall: 0.9930  
F1-score: 0.9877

Class 2:  
Precision: 0.9690  
Recall: 0.9690  
F1-score: 0.9690

Class 3:  
Precision: 0.9696  
Recall: 0.9782  
F1-score: 0.9739

Class 4:  
Precision: 0.9746  
Recall: 0.9756  
F1-score: 0.9751

Class 5:  
Precision: 0.9786  
Recall: 0.9753  
F1-score: 0.9770

Class 6:  
Precision: 0.9810  
Recall: 0.9718  
F1-score: 0.9764

Class 7:  
Precision: 0.9680  
Recall: 0.9698  
F1-score: 0.9689

Class 8:  
Precision: 0.9741  
Recall: 0.9651  
F1-score: 0.9696

Class 9:  
Precision: 0.9749  
Recall: 0.9613  
F1-score: 0.9681

Detailed Classification Report:

	precision	recall	f1-score	support
0	0.98	0.99	0.99	980
1	0.98	0.99	0.99	1135
2	0.97	0.97	0.97	1032
3	0.97	0.98	0.97	1010
4	0.97	0.98	0.98	982
5	0.98	0.98	0.98	892
6	0.98	0.97	0.98	958
7	0.97	0.97	0.97	1028
8	0.97	0.97	0.97	974
9	0.97	0.96	0.97	1009
accuracy			0.98	10000
macro avg	0.98	0.97	0.98	10000
weighted avg	0.98	0.98	0.98	10000

```
In [1]: # Import necessary libraries
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import classification_report

# Create sample data
np.random.seed(42)
n_samples = 300
```

```

# Generate feature data
X = np.random.normal(size=(n_samples, 2))

# Generate target variable (3 classes)
# Class 0: Low values of both features
# Class 1: High values of first feature
# Class 2: High values of second feature
y = np.zeros(n_samples)
y[X[:, 0] > 0.5] = 1
y[X[:, 1] > 1.0] = 2

# Convert to DataFrame for better visualization
df = pd.DataFrame(X, columns=['Feature_1', 'Feature_2'])
df['Target'] = y

# Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Scale the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Train the model
model = LogisticRegression(multi_class='multinomial', solver='lbfgs', max_iter=1000)
model.fit(X_train_scaled, y_train)

# Make predictions
y_pred = model.predict(X_test_scaled)

# Print classification report
print("Classification Report:")
print(classification_report(y_test, y_pred))

# Print coefficients for each class
print("\nModel Coefficients:")
for i, class_label in enumerate(model.classes_):
    print(f"\nClass {class_label}:")
    print(f"Feature 1: {model.coef_[i][0]:.4f}")
    print(f"Feature 2: {model.coef_[i][1]:.4f}")
    print(f"Intercept: {model.intercept_[i]:.4f}")

# Example prediction for new data
new_data = np.array([[0.5, 1.2]])
new_data_scaled = scaler.transform(new_data)
prediction = model.predict(new_data_scaled)
probabilities = model.predict_proba(new_data_scaled)

print("\nPrediction for new data point [0.5, 1.2]:")
print(f"Predicted class: {prediction[0]}")
print("Probabilities for each class:")
for i, prob in enumerate(probabilities[0]):
    print(f"Class {i}: {prob:.4f}")

```

```

Classification Report:
              precision    recall  f1-score   support

     0.0         0.95      0.97      0.96         37
     1.0         0.93      0.93      0.93         15
     2.0         0.86      0.75      0.80          8

 accuracy          0.93         60
 macro avg          0.91      0.89      0.90         60
weighted avg          0.93      0.93      0.93         60

```

#### Model Coefficients:

```

Class 0.0:
Feature 1: -1.9889
Feature 2: -1.7381
Intercept: 2.7600

```

```

Class 1.0:
Feature 1: 2.5937
Feature 2: -1.6915
Intercept: -0.1179

```

```

Class 2.0:
Feature 1: -0.6048
Feature 2: 3.4297
Intercept: -2.6422

```

Prediction for new data point [0.5, 1.2]:

Predicted class: 2.0

Probabilities for each class:

Class 0: 0.1200

Class 1: 0.0792

Class 2: 0.8008

In [5]: probabilities

Out[5]: array([[0.11999349, 0.07923127, 0.80077524]])

In [ ]: